

MARKETING CAMPAIGN ROI & ATTRIBUTION — INSIGHTS MEMO

Prepared by Adham AlHers · Data Analyst / Business Analyst Portfolio Project

Headline finding #1: attribution model choice changes which channel "wins"

SEO ranks #1 in revenue credit under first-touch attribution, but drops to #4 (last) under last-touch. Email shows the exact opposite. Neither model is "correct" on its own — they answer different questions (what starts a journey vs. what closes it).

Channel	First-Touch Revenue	Last-Touch Revenue	Rank Flip
SEO	€284,357 (#1)	€110,536 (#4)	1 → 4
Email	€89,074 (#4)	€275,957 (#1)	4 → 1
Paid Social	€233,750 (#2)	€120,582 (#3)	2 → 3
Paid Search	€110,257 (#3)	€210,363 (#2)	3 → 2

A budget decision based on last-touch data alone would systematically starve SEO — the channel actually bringing new people into the funnel — because SEO rarely gets to "close" a sale it started weeks earlier.

Headline finding #2: Paid Social is over-funded relative to its marginal returns

Paid Social receives 71.0% of total marketing budget but generates only 16.5% of total revenue. More importantly, its ROAS has fallen from 2.61x (early period, lower spend) to 1.94x (late period, higher spend) — a 25.4% decline — as spend has scaled up. This is genuine diminishing returns, not noise: Email, SEO, and Paid Search show no equivalent decline over the same period.

Recommendation

Shift a modest €75,000 from Paid Social to Email. This calculation uses Paid Social's actual marginal ROAS (1.94x — what the next euro spent there is really returning, not its inflated historical average) for what's given up, and a deliberately conservative estimate of Email's incremental return. Projected blended ROAS lift: +17.2%.

Methodology note — read before presenting this number

Email's simulated ROAS (~148x) is unusually high in absolute terms. Rather than take that number at face value when modeling the reallocation, this recommendation deliberately caps Email's assumed incremental ROAS at 50x — roughly a third of its observed average — to account for list-size and deliverability limits that aren't captured in this dataset. A naive extrapolation using the full 148x average would suggest a much larger (and much less credible) lift. Flagging this explicitly is itself part of the analysis: knowing when a number is too good to extrapolate from is what separates a data analyst from someone who just reads a dashboard.